

Ronghui Gu  
51 Prospect Street  
Yale University  
New Haven, CT 06511-8285  
[ronghui.gu@yale.edu](mailto:ronghui.gu@yale.edu)

December 1, 2016

Re: Application for the position of tenure-track assistant professor

Dear Madam/Sir,

Please find enclosed my application materials for the position of tenure-track assistant professor in the Department of Computer Science and Engineering at the University of Minnesota. I am a Ph.D. candidate at Yale University under the direction of Prof. Zhong Shao and will complete my degree in December 2016. I have spoken to Prof. Gopalan Nadathur and he has encouraged me to apply.

My areas of interest are Programming Languages and Operating Systems, more specifically, in techniques for building reliable and secure software systems. My work draws on program verification, certified compilation, formal methods, and language-based support for safety and security.

I am a citizen of China, presently on an F-1 student visa.

I have attached the following documents:

- Curriculum Vitae, including a list of four references (Zhong Shao, Benjamin C. Pierce, Adam Chlipala, and Ruzica Piskac)
- Research Statement
- Teaching Statement
- Three publications in OSDI'16, POPL'15, and PLDI'16

I would be grateful if you could consider my application favorably.

Sincerely,



Ronghui Gu

CONTACT INFORMATION	Department of Computer Science Yale University 51 Prospect Street, New Haven, CT 06511-8937 Phone: +1 (203) 430-2840 Email: ronghui.gu@yale.edu Homepage: <a href="http://www.guronghui.com">http://www.guronghui.com</a>	
INTEREST	Programming languages and operating systems, with a focus on language-based support for safety and security, certified system software, certified programming and compilation, formal methods, and concurrency.	
EDUCATION	<b>Yale University</b>	New Haven, CT
	Ph.D. in Computer Science	2011 - 2016
	<ul style="list-style-type: none"> <li>• Advisor: Prof. Zhong Shao</li> <li>• Thesis: An Extensible Architecture for Building Certified Sequential and Concurrent OS Kernels.</li> <li>• Nomination for ACM Dissertation Award</li> </ul>	
	M.S. and M.Phil. in Computer Science	2011 - 2014
	• GPA: 3.9 / 4	
	<b>Tsinghua University</b>	Beijing, China
	B.S. in Computer Science	2007 - 2011
	<ul style="list-style-type: none"> <li>• Advisor: Prof. Yuan Dong and Prof. Shengyuan Wang</li> <li>• GPA: 91.2 / 100      Rank: 4 / 118</li> <li>• Graduation with Distinction (top 2 %)</li> </ul>	
PROFESSIONAL EXPERIENCE	<b>Yale University</b>	New Haven, CT
	<i>Research Associate</i>	2016 - present
	<i>Research Assistant</i>	2011 - 2016
	Working on CertiKOS, an extensible architecture for building certified OS kernels. As the lead developer of the project, verified a series of sequential and concurrent OS kernels in Coq with my Yale colleagues. The most realistic one is written in 6,500 lines of C and x86 assembly and runs on stock x86 multicore machines.	
	<b>Tsinghua University</b>	Beijing, China
	<i>Research Assistant</i>	Spring 2011
	Working on verifying the preemptive scheduling and nested interrupt handling of $\mu\text{C}/\text{OS-II}$ running on the PowerPC platform.	
CONFERENCE PUBLICATIONS	<b>CertiKOS: An Extensible Architecture for Building Certified Concurrent OS Kernels.</b> (with Z. Shao, H. Chen, X. Wu, J. Kim, V. Sjöberg, and D. Costanzo) <i>Proceedings of the 12th USENIX Symposium on Operating Systems Design and Implementation (OSDI'16), 2016.</i>	

(Accept rate: **18%**)

**End-to-End Verification of Information-Flow Security for C and Assembly Programs.**

(with D. Costanzo and Z. Shao)

*Proceedings of the 37th annual ACM SIGPLAN conference on Programming Language Design and Implementation (PLDI'16), 2016.*

(Accept rate: **16%**)

**Toward Compositional Verification of Interruptible OS Kernels and Device Drivers.**

(with C. Hao, X. Wu, Z. Shao, and J. Lockerman)

*Proceedings of the 37th annual ACM SIGPLAN conference on Programming Language Design and Implementation (PLDI'16), 2016.*

(Accept rate: **16%**)

**Deep Specifications and Certified Abstraction Layers.**

(with J. Koenig, T. Ramanandro, Z. Shao, X. Wu, S. Weng, H. Zhang, and Y. Guo)

*Proceedings of the 42nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL'15), 2015.*

(Accept rate: **23%**)

MANUSCRIPTS     **Language and Compiler Support for Building Certified Concurrent Abstraction Layers.**

(with Z. Shao, X. Wu, J. Kim, J. Koenig, T. Ramanandro, V. Sjoberg, H. Chen, and D. Costanzo)

Submitted to *the 38th annual ACM SIGPLAN conference on Programming Language Design and Implementation (PLDI'17), 2017.*

**Safety and Liveness of MCS LockLayer by Layer.**

(with V. Sjöberg, J. Kim, and Z. Shao)

Submitted to *the 6th ACM SIGPLAN Conference on Certified Programs and Proofs (CPP'17), 2017.*

TEACHING     **Teaching Assistant** at Yale University     Fall 2012 - Fall 2015

EXPERIENCE

- CPSC 422/522 Operating System Design and Implementation (Fall 2015)  
Instructor: Prof. Zhong Shao
- CPSC 458/558 Automatic Decision Systems (Fall 2015)  
Instructor: Dr. Stephen Slade
- CPSC 439/539 Software Engineering (Spring 2015 and Spring 2014)  
Instructor: Prof. Ruzica Piskac
- CPSC 424/524 Parallel Programming Techniques (Fall 2014 and Fall 2013)  
Instructor: Dr. Andrew Sherman
- CPSC 112 Introduction to Programming Languages (Spring 2013 and Fall 2012)  
Instructor: Prof. Drew V. McDermott and Prof. Daniel Abadi

ADVISING EXPERIENCE	<b>Undergraduate students</b> at Yale University	
	<b>Hengchu Zhang</b> , senior project on verifying synchronous IPC	Spring 2015
	<b>Graduate students</b> at Yale University	
	<b>Mengqi Liu</b> , research project on verifying real-time CPS	2015 - present
SELECTED PRESENTATIONS	<b>An Extensible Architecture for Building Certified Concurrent OS Kernels.</b>	
	• University of California, Davis, CA	Dec. 2016
	• Columbia University, NY	Dec. 2016
	• University of Pennsylvania, PA	Nov. 2016
	• Princeton University, NJ	Nov. 2016
	• OSDI, GA	Nov. 2016
	<b>Deep Specifications and Certified Abstraction Layers.</b>	
	• Yale Programming Languages Day, CT	Nov. 2015
	• New England Programming Languages and Systems Symposium	June 2015
	• High Confidence Software and Systems Conference, Maryland	May 2015
	• POPL, Mumbai, India	Jan. 2015
SELECTED HONORS	Nomination for ACM Dissertation Award, Yale University	Aug. 2016
	Robert Willets Carle Scholarship, Yale University	Feb. 2016
	Travel Grant for POPL15, NSF	Jan. 2015
	Doctoral Fellowship, Yale University	Aug. 2011
	Graduation with Distinction (top 2%), Tsinghua University	Jul. 2011
	Outstanding Graduate, Beijing City	Jul. 2011
	Outstanding Undergraduate Thesis Award, Tsinghua University	Jul. 2011
	Sohu Scholarship, Tsinghua University	Oct. 2010
	Ticket Master Scholarship, Tsinghua University	Oct. 2009
	Tsinghua Outstanding Social Work Scholarship	Oct. 2008
	Elite Youth League Member of Jiangsu Province, China	Apr. 2007
	First Prize of National Mathematical Olympiad Competition, China	Oct. 2006
	First Prize of National Chemistry Olympiad Competition, China	Oct. 2006
MEDIA COVERAGE	Yale Daily News	Nov. 2016
	International Business Times	Nov. 2016
	Yale News	Nov. 2016
INDUSTRIAL EXPERIENCE	<b>Google Inc.</b>	Mountain View, CA
	<i>Software Engineer</i>	2016 - present
	<b>Ecopia Tech Corporation</b>	Ontario, Canada
	<i>Research Intern</i>	Summer 2015

PROFESSIONAL  
ACTIVITIES     **Member:** CPS-VO  
                  **Reviewer:** TACAS'15

CITIZENSHIP     Peoples Republic of China

REFERENCES     **Zhong Shao**  
                  Professor  
                  Department of Computer Science  
                  Yale University  
                  51 Prospect Street  
                  New Haven, CT 06511–8937, USA  
                  Email: zhong.shao@yale.edu  
                  Phone: +1 203 432-6828  
                  Homepage: <http://www.cs.yale.edu/homes/shao/>

**Benjamin C. Pierce**  
                  Henry Salvatori Professor  
                  Department of Computer and Information Science  
                  University of Pennsylvania  
                  3330 Walnut Street  
                  Philadelphia, PA 19104-6389, USA  
                  Email: bcpierce@cis.upenn.edu  
                  Phone: +1 215 898-6222  
                  Homepage: <http://www.cis.upenn.edu/~bcpierce/>

**Adam Chlipala**  
                  Associate Professor  
                  Department of Electrical Engineering and Computer Science  
                  MIT  
                  77 Mass Ave, 32-G842  
                  Cambridge, MA 02139–4307, USA  
                  Email: adamc@csail.mit.edu  
                  Phone: +1 617 324-8439  
                  Homepage: <http://adam.chlipala.net/>

**Ruzica Piskac**  
                  Assistant Professor  
                  Department of Computer Science  
                  Yale University  
                  51 Prospect Street  
                  New Haven, CT 06511–8937, USA  
                  Email: ruzica.piskac@yale.edu  
                  Phone: +1 203 432-8001  
                  Homepage: <http://www.cs.yale.edu/homes/piskac/>

# RESEARCH STATEMENT

Ronghui Gu ([ronghui.gu@yale.edu](mailto:ronghui.gu@yale.edu))

My research goal is to make critical software systems truly reliable and secure through *formal verification*. As the backbone of modern software systems, operating system (OS) kernels have great impacts on the reliability and security of today's computing hosts. OS kernels, however, are complicated, highly concurrent, and prone to bugs. For the past several years, my research has focused on investigating compositional programming language theories and developing scalable tools to build verified sequential and concurrent OS kernels that are formally proved to be error-free and secure [POPL'15, OSDI'16, PLDI'16a, PLDI'16b]. Broadly speaking, my research falls into the subfield of programming languages, dealing with the principles and practice of formal verification of real software systems.

## My Current Work on CertiKOS

My research approach is to explore clean and concise programming language theories that reveal the fundamental insights of system design patterns, and to apply these theories in building practical systems that behave as required through formal verification. While such mechanical “formal verification” dates back to the 1960s, complete formal proofs of sequential OS kernels only became feasible recently, as demonstrated by seL4 in 2009. This result was so encouraging that it seemed not too far away from building a fully verified concurrent kernel under reasonable proof efforts. However, seven years have passed, and this last mile is still insurmountable. Even in the single-core setting, the cost of such verification is quite prohibitive (seL4 took 11 person-years to develop), and it is unclear how to quickly adapt such a verified system to support new features and enforce richer properties. Furthermore, none of the previously verified systems have addressed the issues of concurrency.

We believe that these problems are caused by ignoring the *layered structure* in the programming language theories. Although modern systems are implemented with well-designed layers, this layered structure has not been exploited by the verification techniques like program logics: modules across different layers and multiple threads have to be reasoned about at the same abstraction level. This makes system verification difficult to untangle and costly to extend.

My thesis research is among the first to address these challenges by realizing a novel class of specifications, named *deep specifications*, through a layered approach. In theory, these layered deep specifications uncover the insights of layered design patterns and are rich enough to fully characterize a system's functionality. In practice, they are “live” enough to connect with an actual implementation and to provide a modular approach to building entirely trustworthy software system stacks. These advances in both dimensions have resulted in a comprehensive verification framework, named CertiKOS, as well as a series of fully verified sequential and concurrent OS kernels. This CertiKOS work is used in high-profile DARPA programs [CRASH, HACMS], is a core component of an NSF Expeditions in Computing project [DeepSpec], and has been widely considered “a real breakthrough” toward hacker-resistant systems [YaleNews, IBTimes, YDN].

## Deep Specifications and Certified Abstraction Layers

One innovation of my thesis research is that OS kernels are treated as *run-time compilers* [POPL'15]. From this novel view, the OS kernel as a whole compiles user programs that are understood using system call specifications into programs that interact with the kernel implementation directly. We can view a layer in the kernel as a *compilation phase* and the kernel module between two layers as a *compiler transformation*. In this way, OS kernels can be verified by showing that every such transformation preserves the behavior of *arbitrary context programs*, and that all the compilation phases can be composed into a *compositional compiler*. Due to this contextual preservation property, these layer interfaces are named as deep specifications, which precisely capture the functionality of the implementation under any context.

Just as compilers enable program development in higher-level and machine-independent languages, this layer-based approach allows program verification to be done over a more abstract interface that is implementation independent. Each module is verified at its proper abstraction level by showing a contextual simulation between the implementation at that level and the deep specification at a higher level. Once this proof is done, the deep specification forms a new layer interface, which we call a *certified abstraction layer*. Any client program or any functional property of this module can be reasoned about using this more abstract layer specification alone, without looking at the actual implementation.

To apply this layer-based approach to the verification of real software systems, we developed the CertiKOS framework to specify, verify, and compose certified abstraction layers in the Coq proof assistant. We also extended the CompCert verified compiler to compile certified C layers such that they can be linked with assembly layers. With CertiKOS, the task of verifying an OS kernel can be mechanically decomposed into many small, simple, and independent tasks, which are convenient for manual proofs and are often automatable. To demonstrate the power of our framework, we have developed multiple certified sequential OS kernels in Coq. The most realistic one is called mCertiKOS, which consists of 37 certified abstraction layers, took less than one person-year to develop, and can boot a version of Linux as a guest. An extended version of mCertiKOS was deployed on a military land vehicle in the context of a high-profile DARPA program [HACMS].

## Verifying Concurrent OS Kernels with CertiKOS

Moving from the sequential kernel verification to the concurrent one is not straightforward at all and requires a more robust compositional theory. A concurrent kernel allows arbitrarily interleaved execution of kernel/user modules across different layers of abstraction. Several researchers even believe that the combination of concurrency and the kernels' functional complexity makes the formal verification intractable, and even if it is possible, the cost would far exceed that of verifying a single-core sequential kernel.

We believe that the key to untangling this “intractable” complexity is rooted in the strong contextual property of deep specifications [OSDI'16]. A deep specification, in the concurrent setting, must ensure behavior preservation not only under any client context but also under any *concurrent context* with any interleaving. A concurrent context corresponds to a particular execution of the kernel and is represented as a list of *events* that encapsulates the behaviors of the rest of the CPUs (or threads), as well as the interference among them under this execution.

The *certified concurrent layer* is then parameterized over this concurrent context, and a new layer-based theory is developed to compose certified concurrent layers with valid concurrent contexts. Given a particular concurrent context, the interleaving is determined, and a concurrent layer is reduced to a sequential one — this allows us to apply sequential verification techniques for building new concurrent layers. A newly introduced concurrent layer becomes “deep” only after we show that its simulation proof holds for all valid concurrent contexts. To ease this step, we lift the concurrent machine model (which allows arbitrary interleaving at any point) to an abstract local machine model, where all interference from the concurrent context is restricted to certain specific points, i.e., just before the invocations of synchronization primitives. The idea of this machine lifting enables the *local reasoning* of multicore (and multithreaded) programs, without worrying about the interleaving except for a few specific places. Furthermore, based on this machine lifting, we introduce a thread-safe version of the CompCert compiler called CompCertX, which compiles concurrent C layers into assembly ones and propagates the guarantees down to the assembly-level execution over the multicore hardware.

Using this concurrent layer-based theory, we have also developed and verified a practical concurrent OS kernel in Coq. Our certified kernel is written in 6,500 lines of C and x86 assembly, and it runs on stock x86 multicore machines. To our knowledge, this is the first fully verified concurrent OS kernel with fine-grained locking. This work has been recognized as “a milestone that ... could lead to a new generation of reliable and secure systems software” [YaleNews].

## Adapting CertiKOS to Verify Interruptible OS Kernels and Device Drivers

Besides the power to build certified software systems from scratch, our layer-based theory also makes it efficient to introduce new system features. One example is our CertiKOS extensions for verifying device drivers [PLDI'16b]. In a monolithic kernel, device drivers represent the majority of the code base, as well as the primary source of the system crashes. Although formal verification of device drivers is highly desirable, it is widely considered to be challenging, due to the abundance of device features and the non-local effects of interrupts.

To address this issue, we extend the certified layer with both a general device model that can be instantiated with various hardware devices and a realistic interrupt model that scales the reasoning about interruptible code. The device drivers are modeled as if each of them were running on a separate *logical CPU*. This novel idea allows us to incrementally refine a raw device model into a more abstract one by building certified layers of the relevant driver on its logical CPU. Meanwhile, this idea systematically enforces the isolation among different devices and the rest of the kernel. This strong isolation property leads to an abstract interrupt model, under which most of the kernel execution is interrupt-unaware.

Thanks to the layer-based framework equipped with these new models, we turned mCertiKOS into an interruptible kernel with verified device drivers (e.g., serial and IOAPIC). The entire extension took seven person-months to implement in Coq. To the best of our knowledge, this is the first verified interruptible OS kernel with device drivers.

## End-to-End Security Verification in CertiKOS

In the layer-based approach, richer properties of the whole system can be easily derived from the deep specifications. Take the security property as an example. Protecting the confidentiality of information manipulated by a software system is one of the critical challenges facing today's cybersecurity community. It is challenging to specify the desired security policy of a complex system in a clean way and to conduct the proof in a uniform way that allows linking everything together into a system-wide guarantee. We address these issues by extending our CertiKOS framework [PLDI'16a]. We develop a novel methodology, called the *observation function*, which provides a uniform mechanism to specify security policies. A policy can be proved at the top-most layer of the kernel using a general method that subsumes both security-label proofs and information-hiding proofs. By showing that our simulation preserves security, the guarantees can be propagated across layers of abstraction, resulting in the first ever end-to-end secure kernel involving both C and assembly code.

## Future Research Agenda

The goal of my research is to integrate clean and concise programming language theories into the development of formally verified software systems. I plan to pursue this goal through the following research directions at different stages.

### Programming Certified Software Systems Directly

Existing projects on real system verification, including CertiKOS, all require (manually) writing the actual implementation in a C-like language and a formal specification in a proof assistant language. Lacking the support of writing certified programs directly makes the certified software systems difficult to develop and maintain. I plan to work with *program synthesis* and *artificial intelligence* researchers to bridge this enormous gap between the low-level system programming and the high-level specification reasoning through a uniform framework that enables programming certified software directly.

This goal is ambitious but promising. Because deep specifications precisely capture the contextual functionality of their implementations, why not write only the layer specifications and then generate the whole system from the layers automatically? When focusing on a single verification task between two layers, the gap between the implementation and its deep specification, as well as the gap between two adjacent specifications, are both relatively small. We have already synthesized a page allocator module consisting of four layers. I aim to exploit this uniform framework in more depth by attempting: (1) to generate a complete kernel from layer specifications following the line of the program synthesis work; and (2) to link adjacent layers automatically by taking advantage of the recent progress in the artificial intelligence.

### Apply Verification Theories to Various Domains

OS kernels are not the only area that can benefit from the formal verification techniques. In the next five years, I will further investigate more powerful and more compositional verification theories, applying them to the following domains:

- I plan to work with *system researchers* to build a zero-vulnerability system stack consisting of verified components, such as database systems, file systems, network stacks, operating systems, and distributed systems. Although each of these components has been actively studied, for most of them (e.g., file systems), only sequential versions have been verified, and there is a high demand to link all their guarantees together to form a trustworthy system stack. Our success on device driver extensions reveals a promising way to address these challenges. For different systems, we may provide “customized” machine models by exposing an interface that abstracts a particular set of system features. This will enable a domain-specific approach to building certified layers for each system separately. The major remaining issue, then, is how to link together all of these layers with different reliability requirements and system features, using a single compositional theory. I am looking forward to exploring this research opportunity.
- I plan to work with *security researchers* to develop certified commercial-grade toolkits for security protocols (e.g., TLS and SSL). For those layer-based protocols, the layered verification approach will be a perfect fit. The concept of our certified concurrent layers can scale the existing security proofs by instantiating the widely-used oracle techniques with the concurrent context. The theoretical challenge that I have to solve is how to preserve security guarantees under simulations with concurrent contexts.
- I plan to work with *cyber-physical system (CPS) researchers* to build high-confidence CPS with the real-time assurance. Due to the physical consequences, it is highly desirable to prove that a CPS behaves as required regarding both functionality and timing. We have built an embedded system (a variant of mCertiKOS) that is proved to be functionally correct and can run on a real drone. However, the most difficult problem left is how to formally model time in the theory. My proposal is to abstract the flow of time as a list of special events, which will be indicated by a timing oracle that is enriched from the concurrent context. With this model, we can apply our concurrent verification techniques to prove real-time properties. I hope to explore this timing theory shortly.
- I plan to work with *program analysis* researchers to establish a general and efficient specification-based testing framework. For concurrent software (e.g., user-level apps) without high-reliability requirements, the complete formal verification is unnecessary and too expensive, while the random testing has difficulty detecting concurrency bugs caused by certain interleavings. We plan to solve this research problem by utilizing deep specifications to intelligently generate the test cases in the concurrent setting. Due to the “live” property mentioned earlier, we can run these tests over the specifications directly, without waiting for the responses from heavy-workload operations.

### Pushing Towards an Industry-Scale Verification Framework

In the long term, I aim to scale our verification techniques up to the industry grade. To achieve this, we not only have to provide a powerful set of certified libraries but also need to make a breakthrough in the specification theories. I aim to make our deep specifications more natural and straightforward for software engineers. If successful, this research will make the real-world systems reliable and secure, creating a profound impact on the industry and the society in general.



## References

- [POPL'15] R. Gu, J. Koenig, T. Ramanathan, Z. Shao, X. Wu, S. Weng, H. Zhang, and Y. Guo. "Deep specifications and certified abstraction layers." In *42nd ACM Symposium on Principles of Programming Languages (POPL'15)*.
- [OSDI'16] R. Gu, Z. Shao, H. Chen, X. Wu, J. Kim, V. Sjöberg, and D. Costanzo. "CertiKOS: An Extensible Architecture for Building Certified Concurrent OS Kernels." In *12th USENIX Symposium on Operating Systems Design and Implementation (OSDI'16)*.
- [PLDI'16a] D. Costanzo, Z. Shao, and R. Gu. "End-to-end verification of information-flow security for C and assembly programs." In *37th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI'16)*.
- [PLDI'16b] H. Chen, X. Wu, Z. Shao, J. Lockerman, and R. Gu. "Toward compositional verification of interruptible OS kernels and device drivers." In *37th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI'16)*.
- [DeepSpec] DeepSpec: The science of deep specifications. <http://deepspec.org/>.
- [HACMS] High-Assurance Cyber Military Systems (HACMS). <http://opencatalog.darpa.mil/HACMS.html>.
- [CRASH] Clean-slate design of Resilient, Adaptive, Secure Hosts (CRASH). <http://www.darpa.mil/program/clean-slate-design-of-resilient-adaptive-secure-hosts>.
- [YaleNews] "CertiKOS: A breakthrough toward hacker-resistant operating systems." *Yale News*, 2016. <http://news.yale.edu/2016/11/14/certikos-breakthrough-toward-hacker-resistant-operating-systems>.
- [IBTimes] "CertiKOS: Yale develops world's first hacker-resistant operating system." *International Business Times*, 2016. <http://www.ibtimes.co.uk/certikos-yale-develops-worlds-first-hacker-resistant-operating-system-1591712>.
- [YDN] "Yale computer scientists unveil new OS." *Yale Daily News*, 2016. <http://yaledailynews.com/blog/2016/11/18/yale-computer-scientists-unveil-new-os/>.

# TEACHING STATEMENT

Ronghui Gu ([ronghui.gu@yale.edu](mailto:ronghui.gu@yale.edu))

I am excited about teaching because it is not only a procedure to transfer knowledge to students but also a procedure to examine my research and make real impacts by training the next generation of computer scientists.

## Teaching Fellow Experience

My teaching experience spans various activities, such as being a teaching assistant, organizing class projects, and participating in the course development. Among these, I value the teaching assistant role most. During my Ph.D. study at Yale University, I served as teaching fellows for one introductory course (Fall 2012 ~ Spring 2013) and four advanced courses (Fall 2013 ~ Fall 2015). When looking back, it is hard to believe that I have taught more than 600 undergraduate and graduate students directly. These experiences also equip me with the skills to adjust teaching strategies for students with different background and motivations.

Take an introductory course CPSC 112 (Introduction to Programming Languages) as an example. Since most students enrolled in this course were freshmen, and many of them did not have any previous programming experiences, my goal for the class was to inspire their programming interests, instead of emphasizing the coding skills too much. Before the midterm exam, I found that lots of the students were struggling with an assignment related to the regular expression. They were asked to implement some string operations using the regular expression. However, during my office hours, they told me that this concept was too abstract and seemed magical to them. I discussed with the instructor and then gave a tutorial lecture named “Regular Expression in Automata Theory” on the class. I tried to explain the regular expression from the view of automata and utilized slide animations to show how expressions are accepted by the automata. Although such slides took me several days to prepare, it paid when students told me that they finally understood the secrets of regular expression after my lecture.

As for advanced courses like CPSC 424(Parallel Programming Techniques), students are always eager to learn more about debugging and programming skills, which cannot be taught through the lectures alone. To satisfy these demands, I offered a monthly one-to-one training opportunity for all the 15 students in the class. During this training, I worked with the students together on debugging their own parallel code, which is widely considered to be challenging. Sometimes, it took us several hours to locate a single bug. Although it is quite time consuming, I believe this training is very helpful for students to dig out issues in their programs and establish a thorough understanding of how to solve them.

## Course Development Experience

Besides the adaptive teaching approach, I believe that advanced courses are best taught with a combination of theoretical depths and application-oriented practices. With this belief, as the only teaching fellow, I worked with Prof. Ruzica Piskac on developing a new course CPSC 439 (Software Engineering) at Yale University from scratch in Spring 2014. Different from a regular software engineering course that focuses on the engineering side, we insisted on teaching more formal methods. I gave several lectures on how to write specifications, how to reason about code using program logics, and how to use Dafny to assist the development of large projects. All these formal techniques were intensively applied in most of the course projects. At the end of the semester, Ruzica and I held a project fair in the department and invited all the CS faculties, as well as some industry experts from IBM and Microsoft. During the project fair, students presented and demoed their products as if they were running start-ups. Most of the attendants were amazed by the code quality and the reliability-focused design of these projects. Some industry experts said it was unbelievable that these projects were managed by the students under an undergraduate-level course. Furthermore, one project that implemented a secure storing system for medical data has resulted in a start-up.

Another example that demonstrates the power of this combination is our re-design of the course CPSC 422 (Design and Implementation of Operating Systems) at Yale. This re-design was developed by my advisor Prof. Zhong Shao in Fall 2015 based on our CertiKOS research project (i.e., an extensible framework to formally verify OS kernels). As the lead developer of CertiKOS, I actively participated in this course development and gained a new view of the relation between teaching and research. In the past, the OS course at Yale was one of the most challenging courses for undergraduate students. Some students were even scared away by the first code lab when they had a hard time understanding the functionality of the missing code, as well as the role these code play in the whole OS kernel. To help students build a high-level view of the kernel’s design, we established a series of code labs using the layered specifications and implementations of a verified OS kernel called mCertiKOS. We found that this combination of the latest verification results and the OS course not only provided a cleaner view of the kernel concepts and kernel structures, but also made every programming task more precise and concrete. To further carry out this belief, when mentoring students with their course projects, I tried to help them get used to first writing specifications (maybe informal) before starting the actual programming. Some students who

chose to implement device drivers or network stacks told me that these specifications helped them untangle the hardware details and sped up their development. Through this process, two students expressed strong interests in these layered specifications and kept working on their projects after the course.

## Mentoring Experience

Hengchu Zhang is one of the students who continued to work on the OS course project. During his senior project, he attempted to verify an implementation of the synchronous IPC under my guidance. Compared with the original IPC implementation of mCertiKOS, Hengchu's version performs better but heavily depends on the scheduling primitives, which makes it difficult to reason about. During our weekly one-on-one meetings, I helped him reduce this problem into two separate parts: (1) verifying that the implementation satisfies the specification defined in terms of scheduling primitives; and (2) proving the synchronous property over the specification alone. Hengchu solved these problems under my mentorship, and his implementations, along with the proofs, were finally integrated into the CertiKOS code base. After this research experience, Hengchu decided to pursue a Ph.D. degree in the field of formal verification and joined the University of Pennsylvania in Fall 2016.

I enjoyed this mentoring experience a lot. For the past year, I also have been working closely with Mengqi Liu, a second year a Ph.D. student at Yale, on verifying cyber-physical systems with real-time guarantees. Mengqi, our advisor Prof. Zhong Shao, and I divided this project into two phases: (1) porting the existing CertiKOS framework to the ARM platform and verifying the functional correctness of an embedded system; and (2) proving the timing property of the system. I helped Mengqi design the ARM machine model and guided him to finish the first phase. As for the second phase, modeling time in a practical way was a big challenge. Mengqi and I spent several hours a week doing a long run of literature review together and working out a solution to model the flow of time by “counting” the transition steps on the ARM abstract machine. We developed a systematical way to encapsulate these counts into the specification and utilized this information to reason about the timing properties of the whole system. This line of work has been turned into a conference paper submission.

## Courses to Teach

Based on my background, I would like to teach courses in compilers, programming languages, operating systems, formal methods, program verification, and software engineering. I believe that formal methods are best taught in association with sophisticated low-level systems software; and software engineering principles are best conveyed by teaching students programming language theory and foundation of logics. I am looking forward to the opportunity to expanding my teaching and mentoring experience in a world-class university.

**Zhong Shao**

Professor

Department of Computer Science

Yale University

51 Prospect Street

New Haven, CT 06511–8937, USA

Email: [zhong.shao@yale.edu](mailto:zhong.shao@yale.edu)

Phone: +1 203 432-6828

Homepage: <http://www.cs.yale.edu/homes/shao/>

**Benjamin C. Pierce**

Henry Salvatori Professor

Department of Computer and Information Science

University of Pennsylvania

3330 Walnut Street

Philadelphia, PA 19104-6389, USA

Email: [bcpierce@cis.upenn.edu](mailto:bcpierce@cis.upenn.edu)

Phone: +1 215 898-6222

Homepage: <http://www.cis.upenn.edu/~bcpierce/>

**Adam Chlipala**

Associate Professor

Department of Electrical Engineering and Computer Science

MIT

77 Mass Ave, 32-G842

Cambridge, MA 02139–4307, USA

Email: [adamc@csail.mit.edu](mailto:adamc@csail.mit.edu)

Phone: +1 617 324-8439

Homepage: <http://adam.chlipala.net/>

**Ruzica Piskac**

Assistant Professor

Department of Computer Science

Yale University

51 Prospect Street

New Haven, CT 06511–8937, USA

Email: [ruzica.piskac@yale.edu](mailto:ruzica.piskac@yale.edu)

Phone: +1 203 432-8001

Homepage: <http://www.cs.yale.edu/homes/piskac/>